

Zstate Process Mapping - Senior Team Brief

Zstate Equity Research Agent Benchmark
Process Mapping - Senior Team Brief (v0.2)
Prepared for: Zstate.ai | Confidential

1. Purpose

This document provides process mapping, data flows, roles, and timeline for senior stakeholders. It does not include the full 185-task catalog - that is maintained separately for domain experts and task engineers.

2. Strategic Context

Zstate product: Task + Trajectory + Reward datasets for agentic AI training
Benchmark (Phase 1) proves task quality; trajectories (Phase 2) are the product
Differentiation: credentialed experts + real SEC filings + auditable agent runs
MVD: 15 eval tasks, 5 companies, 10-week pilot - not 45 tasks on day one

3. Four-Phase Process Map

Phase 1 - Foundation & Data Corpus

Objective: Single source of truth for filings, transcripts, FX.

Phase 2 - Task Construction

Objective: Expert-authored tasks with ground truth and gold paths.

Phase 3 - Agent Evaluation

Objective: Model-agnostic eval with full trajectory capture.

Phase 4 - Scoring & Release

Objective: Scored benchmark + trajectory export for training.

4. End-to-End Data Flow

SEC EDGAR + Transcripts + FX -> Corpus Service -> Task Registry (prompt + GT)
-> Agent (via adapter) -> Trajectory Logger -> Scoring Engine (L1/L2/L3)
-> Benchmark Record -> Leaderboard + Trajectory Dataset

5. Task Summary (Not Full Catalog)

6. Roles & Capacity

7. Timeline (10 Weeks - MVD)

8. Success Metrics (MVD)

15 published tasks with 100% ground truth citations
>=2 models evaluated with full trajectory capture
>=15pt score spread between best and worst model

>=8 distinct fracture codes documented
Expert k >= 0.7 on 5-task calibration set

9. Decisions Required from Senior Team

Approve 15-task MVD scope (vs 45 on day one)
Confirm expert resourcing (CFA lead + associate)
Approve transcript API budget (~\$100-150/mo for pilot)
Confirm benchmark -> trajectory dataset as product path
Sign off on 10-week pilot timeline

Step | Owner | Input | Output

1.1 CIK mapping | Data Ops | 5 tickers | CIK registry
1.2 EDGAR ingest | Data Ops | 10-K/10-Q | Raw + indexed corpus
1.3 Transcript ingest | Data Ops | Transcript API + IR fallback | Transcript registry
1.4 Section index | Platform Eng | Raw filings | Searchable sections/tables
1.5 Corpus QA | Lead CFA | 10% sample | Signed corpus manifest

Step | Owner | Input | Output

2.1 Template design | Lead CFA | 3 archetypes | Task templates
2.2 Pilot task (GOOGL) | Associate | Corpus | 1 validated task
2.3 Batch authoring | Associate | Templates | 15 task specs
2.4 L1 scripts | Associate + Eng | Ground truth | Verification scripts
2.5 Peer review | Lead CFA | Draft tasks | Published tasks

Step | Owner | Input | Output

3.1 Adapter setup | Platform Eng | Model APIs | Model registry
3.2 Tool sandbox | Platform Eng | Corpus | Search, Parser, Python
3.3 Run campaign | Platform Eng | 15 tasks x N models x 3 | Trajectories
3.4 Fracture analysis | Product + Eng | Trajectories | Fracture report

Step | Owner | Input | Output

4.1 Layer 1 scoring | Platform Eng | Agent output + GT | Hard scores
4.2 Layer 2 scoring | Expert (sample) | Trajectories | Judgment scores
4.3 Layer 3 audit | Platform Eng | Citations | Trust scores
4.4 Release | Product | All artifacts | benchmark_v0.1 + leaderboard
4.5 Trajectory export | Product | Ranked runs | Training dataset v1

Layer | Count | Purpose

Data acquisition | 24 | SEC, IR, peers, FX, citations
Financial statements | 24 | Normalization, footnotes, reconciliation
Financial models | 79 | 3-stmt, DCF, comps, LBO, SOTP, DDM, M&A
Valuation + industry + earnings | 27 | Football field, Porter, guidance drift
Thesis + risk + compliance + output | 31 | Reco, risks, FINRA, initiation memo
Total catalog | 185 | Full analyst workflow indexed
MVD pilot (now) | 15 | 5 cos x 3 archetypes - Type F and M only

Role | Hours/week | Responsibility

Lead CFA | 10-12 | Standards, review, publish approval
MBA Associate | 12-15 | Task authoring, ground truth, citations
Compliance (shared) | 3-5 | FINRA/mandate rules (Type C tasks)
Platform Engineer | Contract/part-time | Corpus, eval runner, L1 scoring

Week | Phase | Milestone

1-2 | 1 + 2 | Corpus for 5 cos; GOOGL pilot task complete
3 | 2 | Templates from pilot; authoring begins
4-6 | 2 + 3 | 15 tasks authored; eval runner live
7-8 | 2 + 4 | Expert review; corpus lock
9-10 | 3 + 4 | Eval campaign; fracture report; release